

SHAURAYA MOHAN

+1 226-989-7032

s35mohan@uwaterloo.ca

Linkedin

Github

shauraya.ca

Education

University of Waterloo

Bachelor of Computer Science, Honours, Co-op - CGPA: 3.9/4.0

January 2025 – Present

Waterloo, ON

Work Experience

Windscribe

January 2026 – August 2026

AI Engineer

Toronto, ON

- Engineered confirmation-gated agent tools for **Garry**, Windscribe's live text-and-voice support agent, enabling autonomous **refunds**, account-detail updates, and changes to user Windscribe settings, cutting human escalations from about **14%** to **2%**.
- Authored the **skill files** and 8 selectable agent **personas** that drive Garry's responses on Windscribe's **Go-based Switchboard** service, keeping tone consistent and on-brand across text and voice.
- Developed **TARS**, an internal **Slack bot** self-hosted on an **Ubuntu VM** that runs a daily **cron** over support tickets through three **GPT-4o** layers (16-category classification, sentiment and churn scoring, QA bug detection), posting digests and a live dashboard for the QA and CS teams.
- Maintained prompts and built **skills** for **OpenClaw**-based internal agents, including an automated **Stripe Radar** review skill that handles all of Windscribe's Stripe Radar fraud reviews.

Control D

May 2026 – August 2026

QA & Automation Engineer

Toronto, ON

- Reproduced and triaged **DNS** filtering issues reported by enterprise users, escalating product bugs to developers and guiding organizations through fixes across **50+** tickets weekly, while managing **SIEM** setups.
- Automated **false-positive/negative** handling for **DNS filter** allowlists and blocklists, auto-verifying each reported case and routing it into the correct list to keep filtering accurate.

ITC

May 2024 – August 2024

Software Development Intern

New Delhi, India

- Built an **Inventory Management System** in **Java** and **SQL** with **real-time audit trails** and automated **shipment tracking**, cutting manual handling by **40%** for a **10-member** logistics team.

Projects

Mogr. - AI-Based Grooming | *Next.js, Supabase, GPT-Vision, MediaPipe*   

- Built 4 independent **GPT-vision** scan pipelines (skin, hair, facial hair, wardrobe) on **Next.js/Supabase**, with a 3-run **self-consistency ensemble** for skin diagnosis and **gpt-image-2** face-preserving previews.
- Architected a 3-stage “**retrieve-then-reason**” outfit recommender (**LLM intent parser** → **deterministic soft-scoring ranker** → **LLM stylist**) that bounds cost by avoiding full-closet LLM reasoning, returning 3 ranked outfits per request.
- Designed a deterministic, **LLM-free color-science engine** (undertone, depth, contrast) and a client-side **MediaPipe Face Mesh** capture gate rejecting unusable photos before they hit a paid vision API.
- Modeled a unified, **RLS-secured Postgres** profile denormalizing scan-derived attributes across 4 feature domains, with request-level **cache-key hashing** to reuse prior LLM results.

Kitchen Copilot - Voice Ordering Assistant | *Next.js, TypeScript, OpenAI Realtime API, Swiggy MCP*  

- Architected a **voice-controlled food-ordering assistant** that places restaurant orders, grocery runs, and table reservations hands-free, built on **WebRTC** and the **OpenAI Realtime API** with server-minted **ephemeral tokens** and low-latency **barge-in** interruption.
- Integrated **Swiggy's food, grocery, and dining platform** through a **streamable-HTTP MCP client** (35 tools, **OAuth 2.1 + PKCE**), with an **LLM intent gateway** routing each request to the right tool set and splitting multi-intent prompts into ordered sub-tasks.

Muse Sketch Studio - Replicate AI Hackathon Winner | *React.js, TypeScript, Node.js, Replicate API* 

- Built an end-to-end **AI fashion design pipeline** (text → sketch → color → model shot → runway video) from prompts using **Gemini nano-banana** and **veo-3** via the **Replicate API**.

Technical Skills

Languages: Python, Java, Go, Ruby, C/C++, JavaScript/TypeScript, SQL, HTML/CSS

AI & ML: LLMs (OpenAI, Anthropic Claude, Gemini), RAG & vector search (Pinecone), prompt engineering, agent orchestration, MCP, computer vision (MediaPipe, TensorFlow, OpenCV), speech (STT/TTS, realtime voice)

Web & Backend: React, Next.js, Node.js, Flask, Django, Tailwind CSS

Tools & Infrastructure: Git, Docker, Linux, Bash, CI/CD, AWS, GCP, PostgreSQL/Supabase, MongoDB, Redis